

The Serverless Shell



Aurèle MAHEO, Pierre SUTRA

Short Bio

- (2015) PhD, High Performance Computing
 - Thesis topic: "Improving the Hybrid model MPI+Threads through Applications, runtimes and Performance tools"
 - Exascale Computing Research



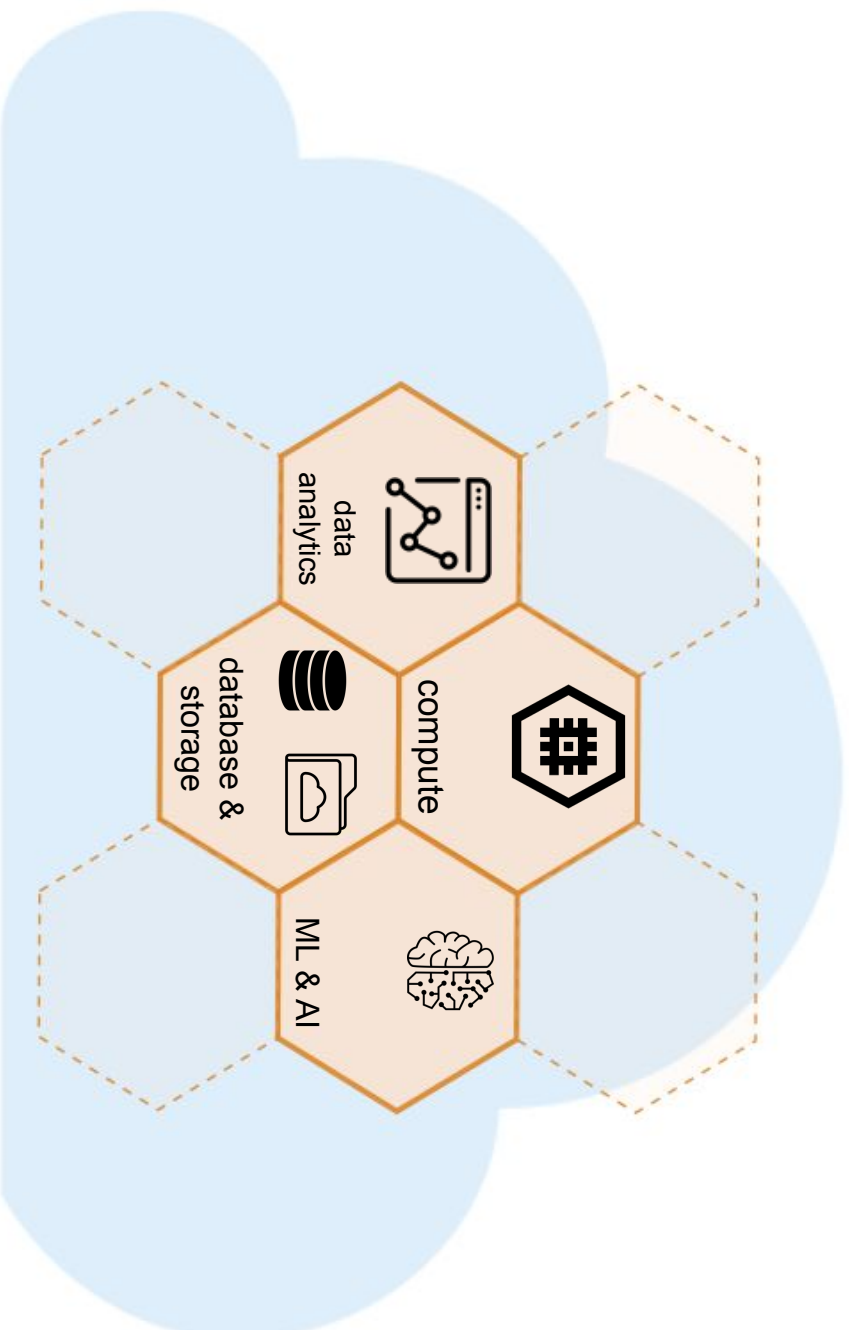
- (2011) MS, High Performance Computing



- (2008) MS, Computer Science




Cloud ecosystem



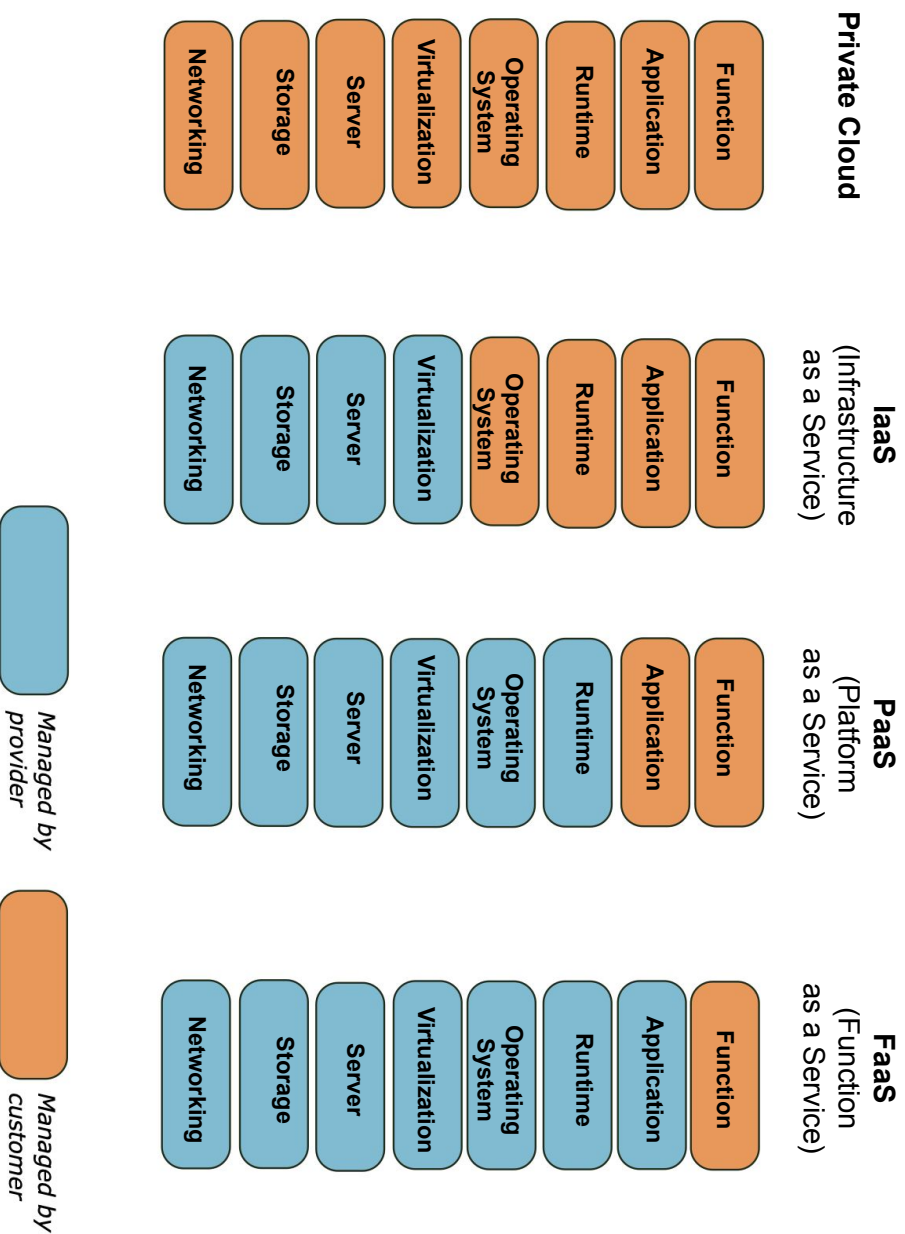
serverless shell

Cloud Computing



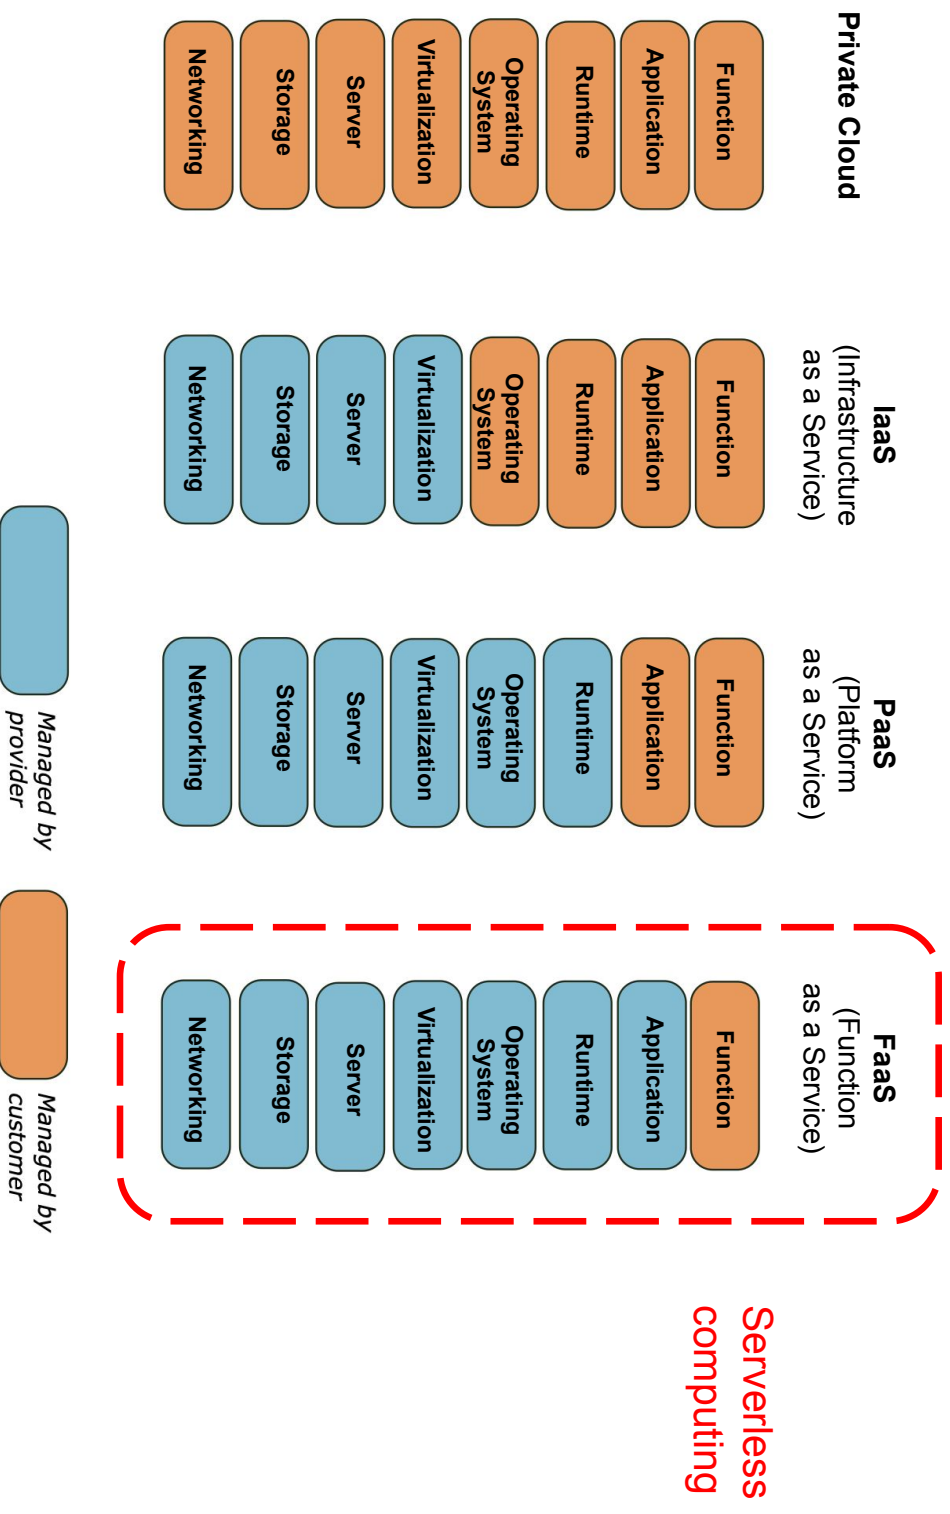
serverless shell

Cloud Computing models



serverless shell


Cloud Computing models



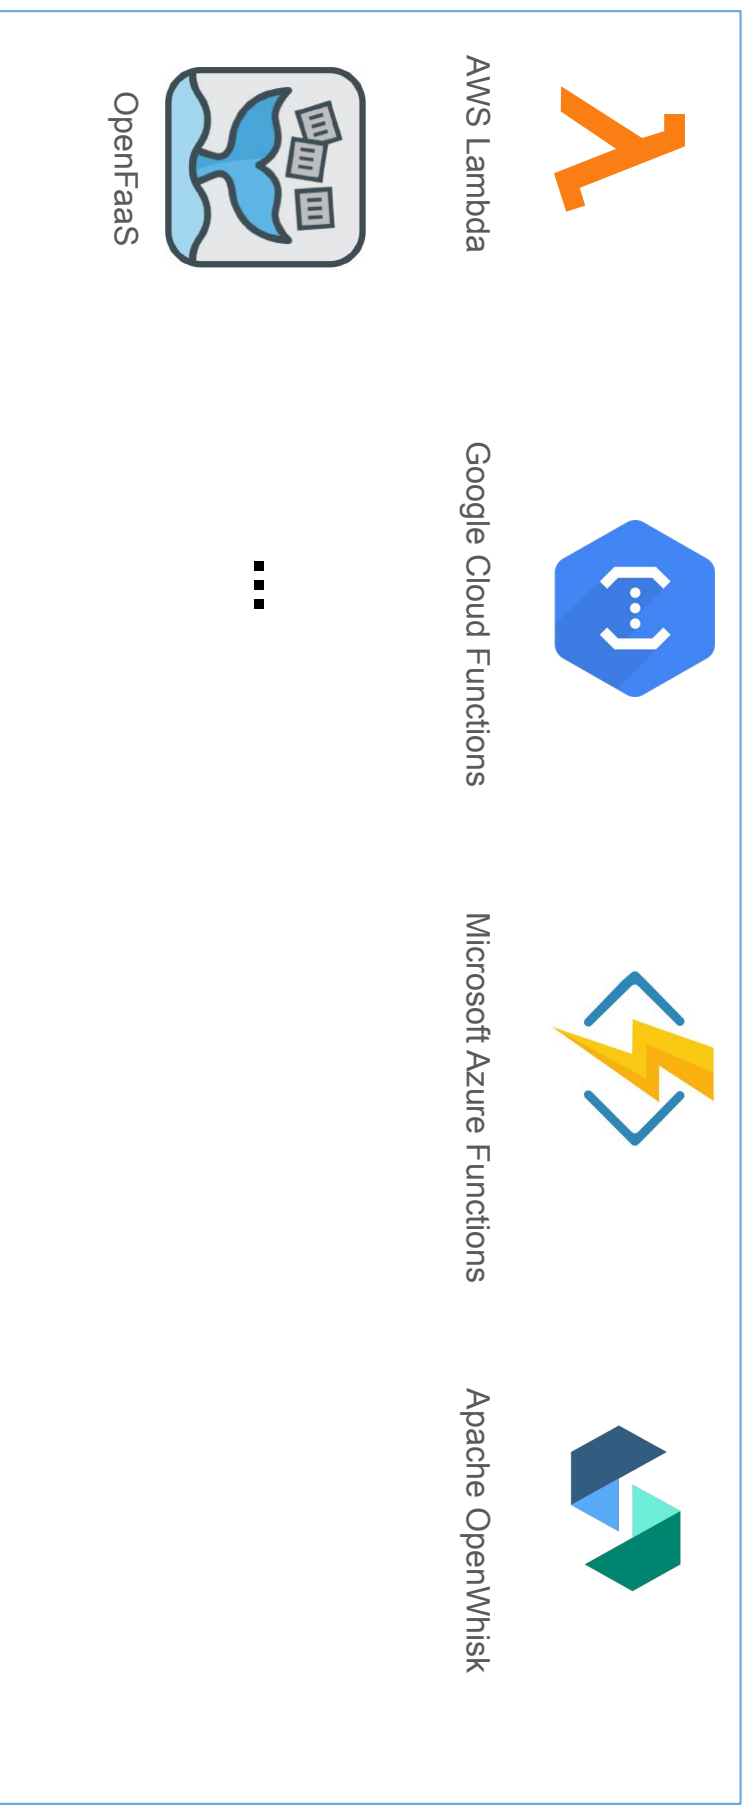
serverless shell

Serverless computing

- Cloud Computing Execution model
- Code execution fully managed by cloud provider
- Seamless server provisioning, administration and maintenance
- Composed of compute runtimes (**F**unction **a**s **a** **S**ervice)




Serverless computing landscape



serverless shell

Example of AWS Lambda



Processing large datasets with UNIX shell

- Languages for data processing
 - Python
 - UNIX shell
 - R
 -

- Why UNIX shell ?

- Semantically minimal language
- Very convenient to process data

```
$> curl dataset | cmdA | cmdB | cmdC | ... | cmdN
```

- Possible to express parallelism using **GNU Parallel**

```
$> cat input | parallel -j<jobs> cmdA
```



serverless shell

Processing large datasets with UNIX shell

- **Limitations**

- Problems arise when dealing with large datasets
- Sequential mode: Long processing
- Parallel mode: Meets bottlenecks ((+100 parallel jobs))
 - Machine is the limit
 - Need of huge local compute resources

- **Solution & Objectives**

- Combine the power of the serverless computing and simplicity of UNIX shell
- Port System mechanisms to the serverless platform
- Adapt UNIX shell for serverless
- Augment it to express stateful patterns

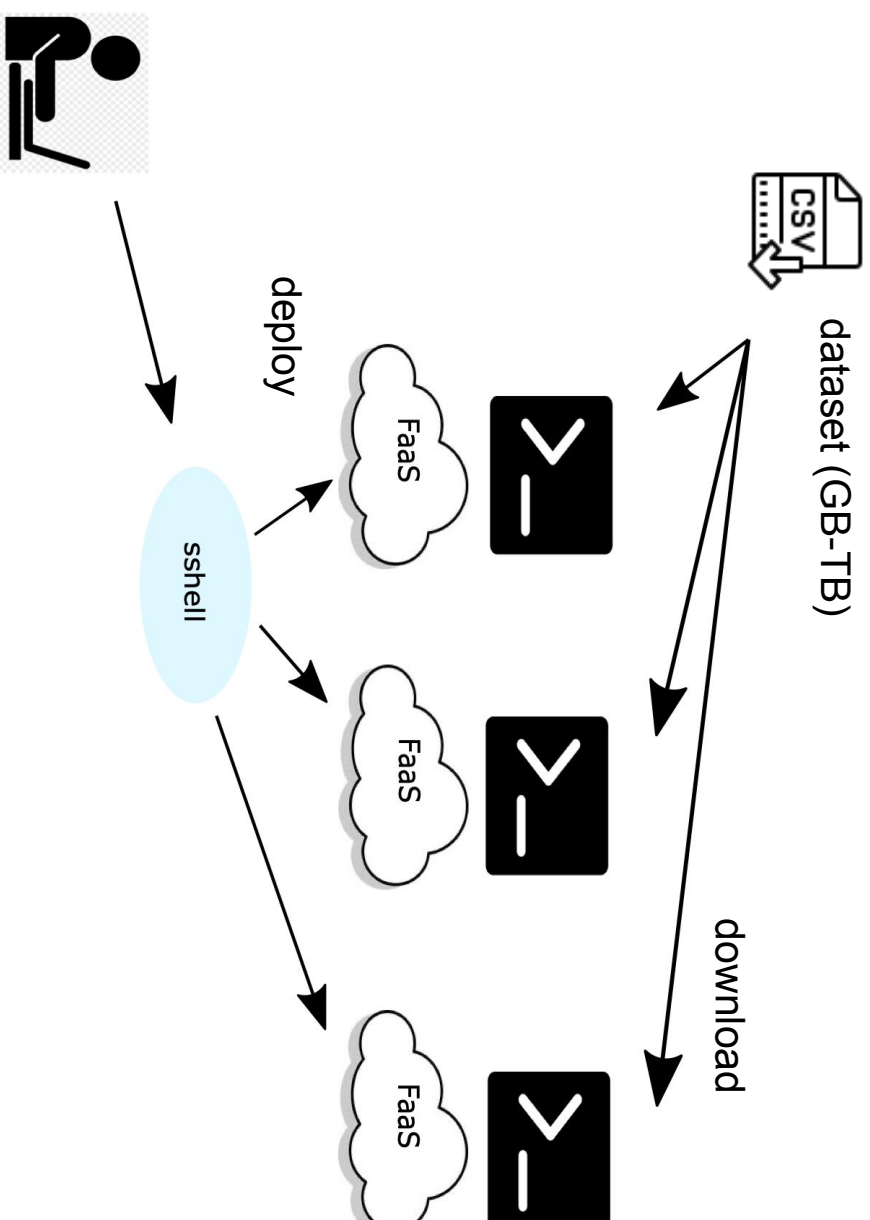
serverless shell

The Serverless Shell

- Usage

```
$> sshell ls -C /  
RequestId: d6215a3a-a41c-4384-b779-215cfa06b30c  
Duration: 13 ms Memory Used: 98 MB  
bin  dev  home  lib64  mnt  proc  run    srv  tmp  var  
boot  etc  lib    media   opt  root  sbin  sys  usr
```

The Serverless Shell



serverless shell

Parallel processing

- Calculate average using **sshell**

```
$> TMP_DIR=/tmp/$(whoami)

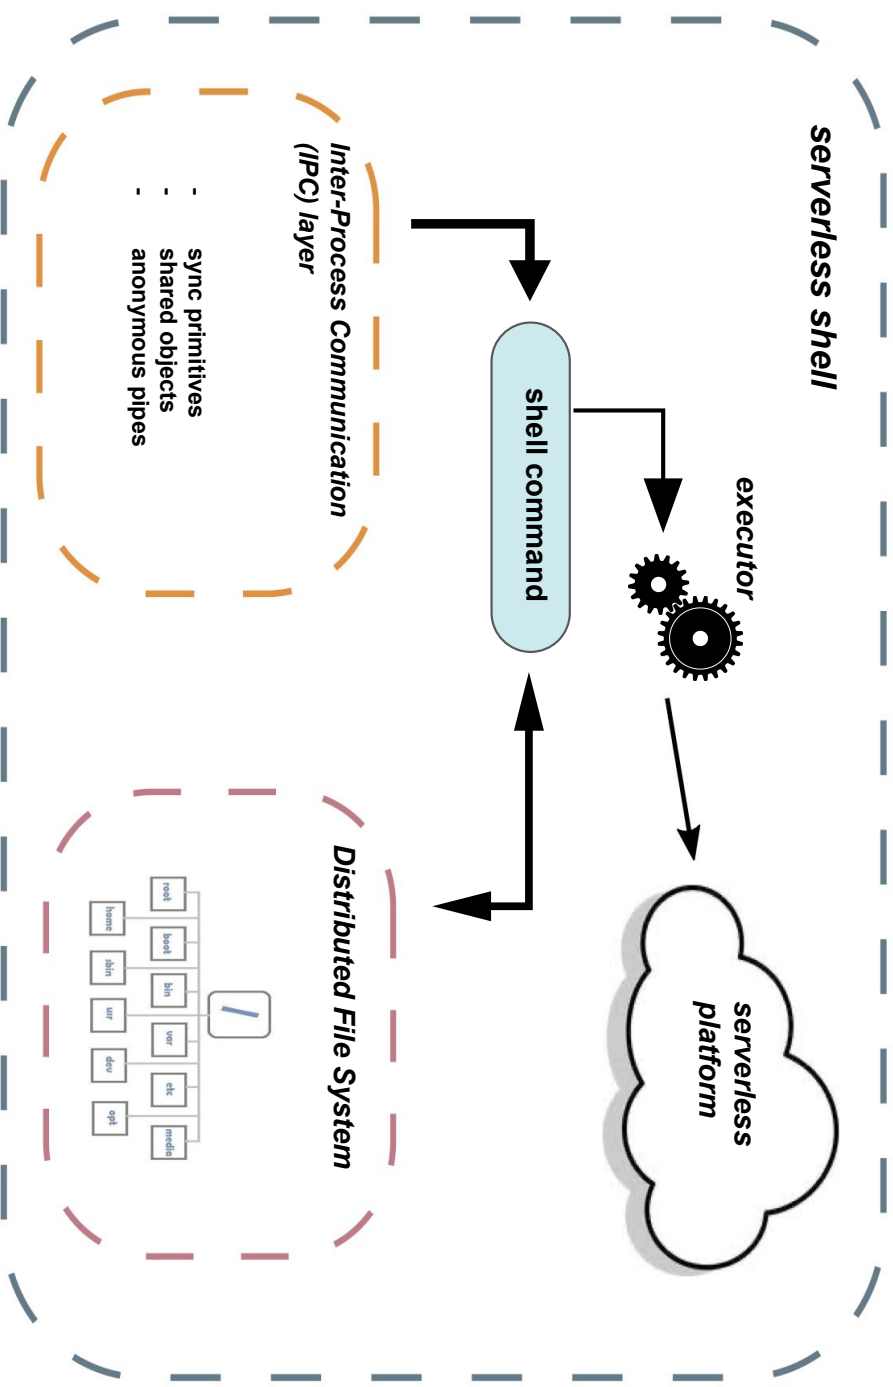
CCBASE="http://commoncrawl.s3.amazonaws.com"
CCMAIN="CC-MAIN-2019-43" # oct. 2019
RANGE="-r 0-10000000"

curl -s ${CCBASE}/crawl-data/${CCMAIN}/warc.paths.gz \
| zcat | head -n ${INPUT} > ${TMP_DIR}/index

average(){
  while read l; do
    sshell "curl -s ${RANGE} ${CCBASE}/${l} | 2>/dev/null zcat -q | grep '^Content-Length ' &
done < ${TMP_DIR}/index | awk '{sum += $2} END {if (NR > 0) print int(sum / NR)}'
```


System design

- Serverless shell:
- serverless platform
 - executor
 - IPC layer
 - Distributed file system



serverless shell

System design / IPC layer



serverless shell

System design / IPC layer

- def: **Stateful** application Versus **Stateless** application

Stateless	Stateful
No use of shared object in application	Use of shared object in application (map, etc)

IPC layer

- sshell example involving synchronization primitive (barrier)

```
$>  
JOBS=100  
BARRIER=$(uid)  
seq 1 1 $((JOBS-1)) | parallel -n0 sshell --async barrier -n  
${BARRIER} await -p ${JOBS}  
sshell barrier -n ${BARRIER} await -p ${JOBS}
```

Inter-Process communication

- Anonymous pipes
- Rewrite “ \$> **cmdA** | **cmdB** ” :

\$>

direction connection

```
sshell "nc -N -l 8080 | cmdB& rdv de41a38e -1 $!P" &  
sshell "HOST=$(rdv de41a38e); exec 3<>/dev/tcp/${HOST}/8080; cmdA >&3; echo EOF  
>&3"
```

\$>

file system

```
sshell "cmdA | awk '{print $0}END{print \"EOF\\\"}' > /fs/de41a38e" &  
sshell "tail -n +0 --pid=$\$ -f --retry /fs/de41a38e 2>/dev/null | { sed \"'/EOF/q'\" && kill \\$\\$ ;} |  
grep -v ^EOF\\$ | cmdB "
```

Implementation

- **Language** : Java
- **SLOC** : ~3K
- **Build system** : Maven
- **Java version** : GraalVM 19.3.0
- **FaaS platform** : AWS Lambda
- **Distributed File system** : AWS Elastic File System (EFS)
- **IPC layer** : Distributed Shared Objects (DSO)

serverless shell


Evaluation

- **Set up**
 - Experiments conducted from an AWS EC2 machine t2.2xlarge (8 vCPUS - 32 GB RAM)
 - Use default parameters in AWS Lambda
 - Each serverless function: 1 GB of memory
- **Performance metrics**
 - Invoke latency
 - I/O (AWS S3 and AWS EFS)
 - Compute
 - Sync
 - Sort

serverless shell


Evaluation / Preliminaries

- sshell built with traditional Java SDK VS Native Java SDK (GraalVM 19.3.0)



serverless shell

Evaluation / Preliminaries



serverless shell

Evaluation / Preliminaries

- Peak transfer rate AWS EFS <> AWS Lambda


	Download	Upload
Sequential	72 MB/s	77 MB/s
Parallel	3418 MB/s	1333 MB/s

Evaluation / Micro Benchmarks

- **Thumbnails generation**
 - Parse a set of 1090 images
 - Generate for each image a 10KB thumbnail
 - Files stored in AWS EFS
- **Port scan analysis**
 - Parse a 40 GB trace containing a full Internet scan of port 80
 - Steps
 - clean raw input data using zannotate
 - isolate Internet Protocol (IP) then Autonomous System (AS)
 - merge the 2 outputs together
 - count the number of IPs and AS

serverless shell

Evaluation / Micro Benchmarks



- **compute component decreases as # parallel jobs increases**
- **invoke latency increases as # parallel jobs increases**

Evaluation / Micro Benchmarks

- Port Scan analysis - native code version

\$>

Step 1 - Annotate

```
cat $JSONFILEEC2 | zannotate -routing -routing-mrt-file=$MRTFILEEC2 -input-file-type=json > $EFSEC2PORTSCANPATH/annotated
```

Step 2 - Extract IP

```
cat $EFSEC2PORTSCANPATH/annotated | jq ".ip" | tr -d '"' > $EFSEC2PORTSCANPATH/extract_ip
```

Step 3 - Extract ASN

```
cat $EFSEC2PORTSCANPATH/annotated | jq -c ".zannotate.routing.asn" > $EFSEC2PORTSCANPATH/extract_asn
```

Step 4 - Calculate popularity

```
pr -mts, $EFSEC2PORTSCANPATH/extract_ip $EFSEC2PORTSCANPATH/extract_asn | awk -F',' '{a[$2]++;} END {for (n in a) print n  
\", \" a[n]} \" | sort -k2 -n -t',' -r > $EFSEC2PORTSCANPATH/as_popularity
```

serverless shell

Evaluation / Micro Benchmarks

- Port Scan analysis - sshell code version

\$>

```
echo Run Port scan analysis - version stateless
JOBS=$1

# Step 1 - split input JSON file into chunks
mkdir $EFSEC2PORTSCANPATH/ckdir
chmod 777 $EFSEC2PORTSCANPATH/ckdir
cd $EFSEC2PORTSCANPATH/ckdir
cat $JSONFILEEC2 | parallel -j200 --pipe --block 40M "cat >
${EFSEC2PORTSCANPATH}/ckdir/cjison_{#}"
split --verbose -n $CHUNKS $JSONFILEEC2 cjison
cd -
durationportscansplitjison=$(expr $clock2 - $clock1)
echo "Port scan 1st part - split: $durationportscansplitjison s"

# Step 2 - annotate each chunk with sshell
mkdir ${EFSEC2PORTSCANPATH}/annotateddir
chmod 777 ${EFSEC2PORTSCANPATH}/annotateddir
ls ${EFSEC2PORTSCANPATH}/ckdir > cklist.out
```

\$>

```
echo size of input elements:
cat cklist.out | wc -l

cat cklist.out | parallel -j$JOBS -l, --env sshell "sshell \' cat
${EFSLAMBDAPORTSCANPATH}/ckdir/., | zannotate -routing
-routing-mrt-file=$MRTFILELAMBDA -input-file-type=json >
$EFSLAMBDAPORTSCANPATH/annotateddir/annotated_${PARALLEL_SEQ}\'""

# Step 3 - parse IP
mkdir $EFSEC2PORTSCANPATH/ipdir
chmod 777 $EFSEC2PORTSCANPATH/ipdir
ls $EFSEC2PORTSCANPATH/annotateddir | parallel -j$JOBS -l, --env sshell
"sshell \' cat ${EFSLAMBDAPORTSCANPATH}/annotateddir/., | jq \'\"\"\".ip\"\"\" >
$EFSLAMBDAPORTSCANPATH/ipdir/ip_${PARALLEL_SEQ}\'\"\""
```

serverless shell

Evaluation / Micro Benchmarks

- Port Scan analysis - sshell code version

\$>

Step 4 - parse ASN

```
mkdir $EFSEC2PORTSCANPATH/asndir
chmod 777 $EFSEC2PORTSCANPATH/asndir
ls $EFSEC2PORTSCANPATH/annotateddir | parallel -j$JOBS -l, --env
sshell "shell \" cat ${EFSLAMBDAPORTSCANPATH}/annotateddir/, | jq
-c \"\"\".zannotate.routing.asn\"\"\" >
$EFSLAMBDAPORTSCANPATH/asndir/asn_${PARALLEL_SEQ}\"\"
#echo $(processaspopularity)
```

Step 5 - Output popularity

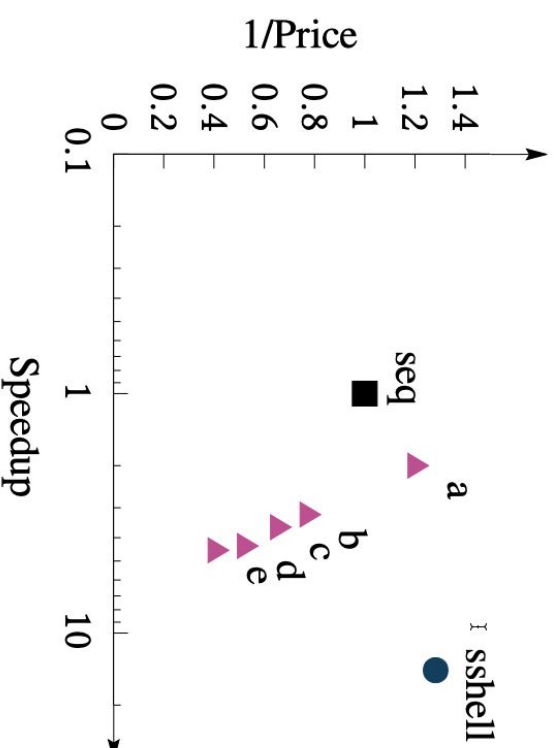
```
cat $EFSEC2PORTSCANPATH/ipdir/ip_* >
$EFSEC2PORTSCANPATH/ipdir/ip_aggr
cat $EFSEC2PORTSCANPATH/asndir/asn_* >
$EFSEC2PORTSCANPATH/asndir/asn_aggr
pr -mts, $EFSEC2PORTSCANPATH/ipdir/ip_aggr
$EFSEC2PORTSCANPATH/asndir/asn_aggr | awk -F' ' '{a[$2]++;} END
{ for (n in a) print n \" \" a[n] } \" | sort -k2 -n -t' ' -r >
$EFSEC2PORTSCANPATH/as_popularity
```

serverless shell

Evaluation / Micro Benchmarks

- Port Scan analysis
 - **seq**: sequential implementation
 - **sshell** : better alternative
 - **a-e** : native execution

ID	c5 EC2 instance (# VCPUs)
a	16
b	36
c	48
d	72
e	98




serverless shell

Evaluation / Large scale application

- **Domain count: Ranking the popularity of web domains**
 - Download archives containing web pages from Commoncrawl (~ 20 TB compressed)
 - Uncompress archives
 - Extracts the outgoing links
 - Count the number of times the domain is mentioned in each page
 - Aggregates results
 - Sort results to construct output

Evaluation / Large scale application




- **compute component decreases as # parallel jobs increases**

serverless shell

Evaluation / Large scale application

- 2 versions for sync component
 - Stateless: use of **AWK** language
 - Stateful: Use of **treemap** DSO object



serverless shell

Evaluation / Large scale application

- LinkRun - Pipeline



serverless shell

Evaluation / Large scale application

- **Sshell version versus LinkRun**

	SLOC	Pricing	Time	Dataset size
Linkrun	716	\$200-260	26-48h	17.62 TB
Sshell	51	\$19	28mn	20.17 TB

serverless shell


Future Work

- PASH : Data-parallel Shell Processing
 - Parallelizes POSIX shell scripts
 - Rewrite them using named pipes (**mkfifo**)
 - Given an input script, enable parallelization using named pipes
- Add back-end to pash:
 - Rewrite PASH output pipes
 - `$> ./pash input.sh -sshell {1,2,3}`

serverless shell

Future Work

- Support multiple FaaS platforms (Google Cloud Platform)
- Connect **sshell** to Kubernetes
 - **Kubernetes** : container orchestrator



References

- "Posh: A Data-Aware Shell", Deepti Raghavan et al. Usenix ATC 2020, Boston, USA. 2020
- "PaSh: light-touch data-parallel shell processing", Nikos Vasilakis et al. EuroSys '21: Proceedings of the Sixteenth European Conference on Computer Systems April 2021

The Serverless shell

